

Can a model transcribe French speech?

A documentation of the model

1 Background

To have feedback for participants in an experiment about clarity, having a model to be as the listener would be a logical next step. Instead of training a model from scratch or even taking a speech to text model and fine tuning it, I figured that there has to be a model that was already trained on French speech.

2 Terminology

I used the term fine tuned, what is it actually. Fine tuning a pre-existing model is just taking a model of choice and further training it on a selected dataset. This makes the model be more of an expert in a certain portion of datatypes. So for example, the model I used that would be utilized is the Huang [2022] whisper-medium model for ASR (automatic speech recognition) in French model which uses OpenAI's medium sized whisper model [Radford et al., 2022, 2023]. The French ASR has been trained "on a composite dataset comprising of over 2200 hours of French speech audio, using the train and the validation splits of Common Voice 11.0, Multilingual LibriSpeech, Voxpopuli, Fleurs, Multilingual TEDx, MediaSpeech, and African Accented French" [Ardila et al., 2020, Pratap et al., 2020, Wang et al., 2021, Conneau et al., 2022, Salesky et al., 2021, Kolobov et al., 2021, gigant, 2022].

ASR models are a type of model that does speech to text transcription. ASR models "converts a speech signal to text, mapping a sequence of audio inputs to text outputs" [Hugging Face, 2024b].

3 Methodology

As mentioned in section 2, I used the whisper-medium model for ASR in French [Huang, 2022]. This is already a powerful model by itself with the amount of hours it was trained on. The model has 769 million parameters to tweak which will come in handy later on.

3.1 Data

Since we have data that was already collected and have the actual transcription, I decided to fine tune the model further with the data. To make sure that the model actually learns to not include any of the annotation anomalies like {np}, {cough}, {sp}, etc, I made sure to remove these. I noticed that if there was some discrepancies with ce and le/la when comes before a vowel so I removed the space that would have been put when extracting from the textgrid. The data can be seen here: lookitsluc1 [2024]. I split the data 80% to 20%. The training split has around 4.59 thousand elements and the test split has around 1.15 thousand elements.

3.2 Fine tuned training

3.2.1 LoRA & PEFT setup

To make the model ready to be fine-tuned, there are some elements that were added on top to get the metrics to a good point. Since whisper-medium-french [Huang, 2022] has 769 million parameters, I did not want to tweak the parameters that gives the power of the whisper model to do the speech-to-text task. To only target the parameters that would best to tweak in my favor, I added a PEFT (Parameter-Efficient Fine-Tuning) via LoRA (Low-Rank Adaptation) wrapper on top of the base model (whisper-medium-french [Huang, 2022]).

LoRA freezes base model's weights and adds an adapter on top that runs along side model and is task specific. PEFT is a way to lower computation times and tunes a small subset of parameters for the task. PEFT takes additional parameters and changes and tweaks them. Adds small, trainable modules or updates a limited set of parameters instead of the full model; Learns task-specific patterns without affecting the core pre-trained model; Maintains efficiency while enabling easy customization using techniques like adapters, LoRA, and prompt tuning [GeeksforGeeks, 2023].

To setup the LoRA and PEFT for the whisper-medium-french [Huang, 2022], there are some parameters I explicitly set. The first is rank (r), which is the amount of parameters to be trained so "a higher rank means the model has more parameters to train, but it also means the model has more learning capacity," [Hugging Face, 2024a]. In the model I have set the rank to 8. The lora_alpha is the scaling factor which controls how intense the adapter weights affect the base model; I put as 16. The bias is "whether none, all or only the LoRA bias parameters should be trained" [Hugging Face, 2024a] which I set to "none". The target_modules are there to let LoRA/PEFT what to focus on when LoRA updates the matrices [Hugging Face, 2023]. Task_type is what the model should be trained for, since I already specified in the loading of the model to be transcribe, I set this to none. Finally lora_dropout randomly turns off a set percentage of adapter weights during training.

3.2.2 Training setup

When it comes to the training arguments, the basic thing to elaborate is that the model ran for 2 epochs, with a learning rate of $2e^{-4}$, a warmup rate of 0.1. There is a weight decay

of 0.05, but this will not need to be worried. Since this is a transcription model, I have set the scale of greediness to the max (1 is greedy in this case). The generation max length that the model is allowed is 200 which means that the model can only generate an output of 200 tokens. To deal with the gradients from peaking, I have set the normalization to 1 (a scale from 0 to 1).

The base model has none of the setup, it is just loaded from the repository and tested on the dataset.

3.2.3 Worry about fine-tuning the model

The main concern with fine tuning the model with the data we have is that the model would memorize the output. Since the transcription are in the same style, the model has a chance of seeing this and just “hack” the system and provide poor metrics. The other worry is with the size of the dataset and how short the audio are, there could be a chance of overfitting, where the model learns the outputs a bit too well. Overfitting can be confusing with the memorization. With both, the main sign is that the model will do extremely well in the training split but when it comes to the test split, there is a gap in the metrics.

When testing the right parameters, I did notice the memorization effect on the model. This could well be on the dataset side plus the extreme tweaking of the model. Overall I got the model to have an acceptable range for French.

If we would use the fine-tuned model, the third worry is that the model is confident in the acoustic of the room and the setup of the mics, so if there are any changes to the room, microphones, background noise, etc, the model might hallucinate and/or overcompensate the output.

4 Results

4.1 Metrics

The metrics used for ASR are Character Error Rate (CER) and Word Error Rate (WER) [Jurafsky and Martin, 2023, Poniro, 2022, Hugging Face, 2022]. These two are the standard. For the fine-tuned model, I kept track on the loss (normal loss and evaluation loss) along with the gradient normalization.

CER is just how good the model can do with transcription [Tamanna, 2025], this is mainly character by character. WER is how many words the model got right [NLPlanet, 2022]. The lower the value of WER, the more accurate the model is. For CER, it is the same, the acceptable range is 5% to 10% and 5% to 15% for WER.

4.2 Base model

The model without any finetuning perform well with the full dataset (train and test split combined). The CER for the model was around 4.5% and a WER of 15.5%. This is good for what the experiment we are doing.

4.3 Fine tuned model

References

Rosana Ardila, Megan Branson, Tim Davis, Michael Henretty, Michael Kohler, Josh Meyer, Reuben Morais, Lindsay Saunders, Francis M. Tyers, and Gregor Weber. Common voice: A massively-multilingual speech corpus. In *Proceedings of the 12th Language Resources and Evaluation Conference*, pages 4218–4222, 2020.

Alexis Conneau, Min Ma, Simran Khanuja, Yu Zhang, Vera Axelrod, Siddharth Dalmia, Jason Riesa, Clara Rivera, and Ankur Bapna. Fleurs: Few-shot learning evaluation of universal representations of speech. *arXiv preprint arXiv:2205.12446*, 2022. URL <https://arxiv.org/abs/2205.12446>.

GeeksforGeeks. What is parameter-efficient fine-tuning (PEFT)? <https://www.geeksforgeeks.org/artificial-intelligence/what-is-parameter-efficient-fine-tuning-peft/>, 2023. Accessed: 2026-08-12.

gigant. African accented french dataset, 2022. URL https://huggingface.co/datasets/gigant/african_accented_french.

Bofeng Huang. Fine-tuned whisper-medium model for asr in French, 2022. URL <https://huggingface.co/bofenghuang/whisper-medium-french>.

Hugging Face. Evaluate: A library for evaluating machine learning models and datasets. <https://github.com/huggingface/evaluate>, 2022.

Hugging Face. PEFT conceptual guide: LoRA. https://huggingface.co/docs/peft/main/en/conceptual_guides/lora, 2023. Accessed: 2026-08-12.

Hugging Face. PEFT documentation: LoRA usage (v0.20.0). https://huggingface.co/docs/peft/v0.20.0/en/package_reference/lora#usage, 2024a. Accessed: 2026-08-12.

Hugging Face. Automatic speech recognition task guide – transformers documentation. <https://huggingface.co/docs/transformers/en/tasks/asr>, 2024b. Accessed: 2026-08-12.

Daniel Jurafsky and James H. Martin. *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*. Pearson, 3rd draft edition, 2023.

Rostislav Kolobov, Olga Okhupkina, Andrey Platunov Olga Omelchishina, Roman Bedyakin, Vyacheslav Moshkin, Dmitry Menshikov, and Nikolay Mikhaylovskiy. Mediaspeech: Multilanguage asr benchmark and dataset, 2021.

lookitsluc1. French cleer speech. https://huggingface.co/datasets/lookitsluc1/french_cleer_speech, 2024.

- NLPlanet. Two minutes NLP: Intro to word error rate (WER) for speech-to-text. <https://medium.com/nlplanet/two-minutes-nlp-intro-to-word-error-rate-wer-for-speech-to-text-fc17a98003ea>, 2022. Accessed: 2026-08-12.
- Nikolas Poniros. jiwer: A Python package for calculating WER, CER, and MER. <https://github.com/jitsi/jiwer>, 2022.
- Vineel Pratap, Qiantong Xu, Anuroop Sriram, Gabriel Synnaeve, and Ronan Collobert. Mls: A large-scale multilingual dataset for speech research. *ArXiv*, abs/2012.03411, 2020.
- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision, 2022. URL <https://arxiv.org/abs/2212.04356>.
- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *International Conference on Machine Learning*, pages 28492–28518. PMLR, 2023.
- Elizabeth Salesky, Matthew Wiesner, Jacob Bremerman, Roldano Cattoni, Matteo Negri, Marco Turchi, Douglas W. Oard, and Matt Post. Multilingual tedx corpus for speech recognition and translation. In *Proceedings of Interspeech*, 2021.
- Tamanna. Deciphering accuracy: Evaluation metrics in NLP and OCR – A comparison of character error rate (CER) and word error rate (WER). <https://medium.com/@tam.tamanna18/deciphering-accuracy-evaluation-metrics-in-nlp-and-ocr-a-comparison-of-character-error-rate-cer-e97e809be0c8>, 2025. Accessed: 2026-08-12.
- Changhan Wang, Morgane Riviere, Ann Lee, Anne Wu, Chaitanya Talnikar, Daniel Haziza, Mary Williamson, Juan Pino, and Emmanuel Dupoux. VoxPopuli: A large-scale multilingual speech corpus for representation learning, semi-supervised learning and interpretation. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 993–1003, Online, August 2021. Association for Computational Linguistics. URL <https://aclanthology.org/2021.acl-long.80>.